

# dlnd\_face\_generation

March 31, 2020

## 1 Face Generation

In this project, you'll define and train a DCGAN on a dataset of faces. Your goal is to get a generator network to generate *new* images of faces that look as realistic as possible!

The project will be broken down into a series of tasks from **loading in data to defining and training adversarial networks**. At the end of the notebook, you'll be able to visualize the results of your trained Generator to see how it performs; your generated samples should look like fairly realistic faces with small amounts of noise.

### 1.0.1 Get the Data

You'll be using the [CelebFaces Attributes Dataset \(CelebA\)](#) to train your adversarial networks.

This dataset is more complex than the number datasets (like MNIST or SVHN) you've been working with, and so, you should prepare to define deeper networks and train them for a longer time to get good results. It is suggested that you utilize a GPU for training.

### 1.0.2 Pre-processed Data

Since the project's main focus is on building the GANs, we've done *some* of the pre-processing for you. Each of the CelebA images has been cropped to remove parts of the image that don't include a face, then resized down to 64x64x3 NumPy images. Some sample data is show below.

If you are working locally, you can download this data [by clicking here](#)

This is a zip file that you'll need to extract in the home directory of this notebook for further loading and processing. After extracting the data, you should be left with a directory of data `processed_celeba_small/`

```
In [1]: # can comment out after executing
        !unzip processed_celeba_small.zip
```

```
Archive:  processed_celeba_small.zip
replace processed_celeba_small/.DS_Store? [y]es, [n]o, [A]ll, [N]one, [r]ename: ^C
```

```
In [40]: data_dir = 'processed_celeba_small/'
```

```
"""
```

```

DON'T MODIFY ANYTHING IN THIS CELL
"""
import pickle as pkl
import matplotlib.pyplot as plt
import numpy as np
import problem_unittests as tests
#import helper

%matplotlib inline

```

## 1.1 Visualize the CelebA Data

The [CelebA](#) dataset contains over 200,000 celebrity images with annotations. Since you're going to be generating faces, you won't need the annotations, you'll only need the images. Note that these are color images with 3 color channels (RGB) each.

### 1.1.1 Pre-process and Load the Data

Since the project's main focus is on building the GANs, we've done *some* of the pre-processing for you. Each of the CelebA images has been cropped to remove parts of the image that don't include a face, then resized down to 64x64x3 NumPy images. This *pre-processed* dataset is a smaller subset of the very large CelebA data.

There are a few other steps that you'll need to **transform** this data and create a **DataLoader**.

**Exercise: Complete the following `get_dataloader` function, such that it satisfies these requirements:**

- Your images should be square, Tensor images of size `image_size x image_size` in the x and y dimension.
- Your function should return a `DataLoader` that shuffles and batches these Tensor images.

**ImageFolder** To create a dataset given a directory of images, it's recommended that you use PyTorch's [ImageFolder](#) wrapper, with a root directory `processed_celeba_small/` and data transformation passed in.

```

In [41]: # necessary imports
import torch
from torchvision import datasets
from torchvision import transforms

In [42]: def get_dataloader(batch_size, image_size, data_dir='processed_celeba_small/'):
    """
    Batch the neural network data using DataLoader
    :param batch_size: The size of each batch; the number of images in a batch
    :param img_size: The square size of the image data (x, y)
    :param data_dir: Directory where image data is located
    :return: DataLoader with batched data
    """

```

```

"""

# TODO: Implement function and return a dataloader
image_aug = transforms.Compose([transforms.Resize(image_size), transforms.CenterCrop(image_size)])
imagenet_data = datasets.ImageFolder(data_dir, transform=image_aug)
data_loader = torch.utils.data.DataLoader(imagenet_data,
                                          batch_size=batch_size + 1,
                                          shuffle=True)

return data_loader

```

## 1.2 Create a DataLoader

**Exercise: Create a DataLoader** `celeba_train_loader` with appropriate hyperparameters. Call the above function and create a dataloader to view images. \* You can decide on any reasonable `batch_size` parameter \* Your `image_size` **must be** 32. Resizing the data to a smaller size will make for faster training, while still creating convincing images of faces!

```

In [43]: # Define function hyperparameters
        batch_size = 32
        img_size = 32

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

# Call your function and get a dataloader
celeba_train_loader = get_dataloader(batch_size, img_size)

```

Next, you can view some images! You should see square images of somewhat-centered faces.

Note: You'll need to convert the Tensor images into a NumPy type and transpose the dimensions to correctly display an image, suggested `imshow` code is below, but it may not be perfect.

```

In [5]: # helper display function
def imshow(img):
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))


"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

# obtain one batch of training images
dataiter = iter(celeba_train_loader)
images, _ = dataiter.next() # _ for no labels

# plot the images in the batch, along with the corresponding labels
fig = plt.figure(figsize=(20, 4))
plot_size=20
for idx in np.arange(plot_size):

```

```
ax = fig.add_subplot(2, plot_size/2, idx+1, xticks=[], yticks=[])
imshow(images[idx])
```



**Exercise: Pre-process your image data and scale it to a pixel range of -1 to 1** You need to do a bit of pre-processing; you know that the output of a tanh activated generator will contain pixel values in a range from -1 to 1, and so, we need to rescale our training images to a range of -1 to 1. (Right now, they are in a range from 0-1.)

```
In [6]: # TODO: Complete the scale function
def scale(x, feature_range=(-1, 1)):
    ''' Scale takes in an image x and returns that image, scaled
        with a feature_range of pixel values from -1 to 1.
        This function assumes that the input x is already scaled from 0-1. '''
    # assume x is scaled to (0, 1)
    # scale to feature_range and return scaled x
    min, max = feature_range
    x = x * (max - min) + min
    return x
```

```
In [7]: """
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

# check scaled range
# should be close to -1 to 1
img = images[0]
scaled_img = scale(img)

print('Min: ', scaled_img.min())
print('Max: ', scaled_img.max())
```

```
Min:  tensor(-1.)
Max:  tensor(0.7176)
```

## 2 Define the Model

A GAN is comprised of two adversarial networks, a discriminator and a generator.

## 2.1 Discriminator

Your first task will be to define the discriminator. This is a convolutional classifier like you've built before, only without any maxpooling layers. To deal with this complex data, it's suggested you use a deep network with **normalization**. You are also allowed to create any helper functions that may be useful.

### Exercise: Complete the Discriminator class

- The inputs to the discriminator are 32x32x3 tensor images
- The output should be a single value that will indicate whether a given image is real or fake

```
In [8]: import torch.nn as nn
import torch.nn.functional as F

In [9]: # helper conv function
def conv(in_channels, out_channels, kernel_size, stride=2, padding=1, batch_norm=True):
    """Creates a convolutional layer, with optional batch normalization.
    """
    layers = []
    conv_layer = nn.Conv2d(in_channels, out_channels,
                           kernel_size, stride, padding, bias=False)

    # append conv layer
    layers.append(conv_layer)

    if batch_norm:
        # append batchnorm layer
        layers.append(nn.BatchNorm2d(out_channels))

    # using Sequential container
    return nn.Sequential(*layers)

In [10]: class Discriminator(nn.Module):
    def __init__(self, conv_dim):
        """
        Initialize the Discriminator Module
        :param conv_dim: The depth of the first convolutional layer
        """
        super(Discriminator, self).__init__()
        self.conv_dim = conv_dim
        # complete init function
        self.conv1 = conv(3, conv_dim, 4, batch_norm=False) # x, y = 64 depth = 3
        self.conv2 = conv(conv_dim, conv_dim * 2, 4) # x, y = 32 depth = 64
        self.conv3 = conv(conv_dim * 2, conv_dim * 4, 4) # x, y = 16 depth = 128

        self.fc = nn.Linear(conv_dim*4*4*4, 1)
        self.out = nn.Sigmoid()
        self.dropout = nn.Dropout(0.5)
```

```

def forward(self, x):
    """
    Forward propagation of the neural network
    :param x: The input to the neural network
    :return: Discriminator logits; the output of the neural network
    """
    # define feedforward behavior
    x = F.leaky_relu(self.conv1(x), 0.2)
    # x = self.dropout(x)
    x = F.leaky_relu(self.conv2(x), 0.2)
    # x = self.dropout(x)
    x = F.leaky_relu(self.conv3(x), 0.2)
    # x = self.dropout(x)

    x = x.view(-1, self.conv_dim*4*4*4)

    x = self.fc(x)
    x = self.dropout(x)

    return x

    """
    DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
    """
    tests.test_discriminator(Discriminator)

```

Tests Passed

## 2.2 Generator

The generator should upsample an input and generate a *new* image of the same size as our training data 32x32x3. This should be mostly transpose convolutional layers with normalization applied to the outputs.

### Exercise: Complete the Generator class

- The inputs to the generator are vectors of some length `z_size`
- The output should be a image of shape 32x32x3

```

In [11]: # helper deconv function
def deconv(in_channels, out_channels, kernel_size, stride=2, padding=1, batch_norm=True):
    """Creates a transpose convolutional layer, with optional batch normalization.
    """
    layers = []
    # append transpose conv layer

```

```

layers.append(nn.ConvTranspose2d(in_channels, out_channels, kernel_size, stride, padding=padding))
# optional batch norm layer
if batch_norm:
    layers.append(nn.BatchNorm2d(out_channels))
return nn.Sequential(*layers)

```

In [12]: `class Generator(nn.Module):`

```

def __init__(self, z_size, conv_dim):
    """
    Initialize the Generator Module
    :param z_size: The length of the input latent vector, z
    :param conv_dim: The depth of the inputs to the *last* transpose convolutional
    """
    super(Generator, self).__init__()

    # complete init function
    self.conv_dim = conv_dim

    self.fc = nn.Linear(z_size, conv_dim*4*4*4)

    self.t_conv1 = deconv(conv_dim*4, conv_dim*2, 4)
    self.t_conv2 = deconv(conv_dim*2, conv_dim, 4)
    self.t_conv3 = deconv(conv_dim, 3, 4, batch_norm=False)
    self.dropout = nn.Dropout(0.5)

def forward(self, x):
    """
    Forward propagation of the neural network
    :param x: The input to the neural network
    :return: A 32x32x3 Tensor image as output
    """
    # define feedforward behavior
    x = self.fc(x)
    x = self.dropout(x)

    x = x.view(-1, self.conv_dim*4, 4, 4)

    x = F.relu(self.t_conv1(x))
    # x = self.dropout(x)
    x = F.relu(self.t_conv2(x))
    # x = self.dropout(x)
    x = F.tanh(self.t_conv3(x))
    return x

    """
    DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
    """

```

```
tests.test_generator(Generator)
```

Tests Passed

## 2.3 Initialize the weights of your networks

To help your models converge, you should initialize the weights of the convolutional and linear layers in your model. From reading the [original DCGAN paper](#), they say: > All weights were initialized from a zero-centered Normal distribution with standard deviation 0.02.

So, your next task will be to define a weight initialization function that does just this!

You can refer back to the lesson on weight initialization or even consult existing model code, such as that from [the networks.py file in CycleGAN Github repository](#) to help you complete this function.

### Exercise: Complete the weight initialization function

- This should initialize only **convolutional** and **linear** layers
- Initialize the weights to a normal distribution, centered around 0, with a standard deviation of 0.02.
- The bias terms, if they exist, may be left alone or set to 0.

```
In [13]: from torch.nn import init
def weights_init_normal(m):
    """
    Applies initial weights to certain layers in a model .
    The weights are taken from a normal distribution
    with mean = 0, std dev = 0.02.
    :param m: A module or layer in a network
    """
    init_gain=0.02
    # classname will be something like:
    # 'Conv', 'BatchNorm2d', 'Linear', etc.
    classname = m.__class__.__name__

    # TODO: Apply initial weights to convolutional and linear layers

    # classname will be something like:
    # 'Conv', 'BatchNorm2d', 'Linear', etc.
    classname = m.__class__.__name__
    # print(classname)
    if hasattr(m, 'weight') and (classname.find('Conv') != -1 or classname.find('Linear') != -1):
        init.normal_(m.weight.data, 0.0, init_gain)
        if hasattr(m, 'bias') and m.bias is not None:
            init.constant_(m.bias.data, 0.0)
    elif classname.find('BatchNorm2d') != -1: # BatchNorm Layer's weight is not a matrix
        init.normal_(m.weight.data, 1.0, init_gain)
        init.constant_(m.bias.data, 0.0)
```



## 2.4 Build complete network

Define your models' hyperparameters and instantiate the discriminator and generator from the classes defined above. Make sure you've passed in the correct input arguments.

```
In [44]: """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        def build_network(d_conv_dim, g_conv_dim, z_size):
            # define discriminator and generator
            D = Discriminator(d_conv_dim)
            G = Generator(z_size=z_size, conv_dim=g_conv_dim)

            # initialize model weights
            D.apply(weights_init_normal)
            G.apply(weights_init_normal)

            print(D)
            print()
            print(G)

            return D, G
```

### Exercise: Define model hyperparameters

```
In [37]: # Define model hyperparams
        d_conv_dim = 64
        g_conv_dim = 64
        z_size = 100

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """

        D, G = build_network(d_conv_dim, g_conv_dim, z_size)
```

```
Discriminator(
  (conv1): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
  )
  (conv2): Sequential(
    (0): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (conv3): Sequential(
    (0): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (fc): Linear(in_features=4096, out_features=1, bias=True)
  (out): Sigmoid()
```

```

        (dropout): Dropout(p=0.5)
    )

    Generator(
      (fc): Linear(in_features=100, out_features=4096, bias=True)
      (t_conv1): Sequential(
        (0): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
      (t_conv2): Sequential(
        (0): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
        (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
      (t_conv3): Sequential(
        (0): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      )
      (dropout): Dropout(p=0.5)
    )

```

### 2.4.1 Training on GPU

Check if you can train on GPU. Here, we'll set this as a boolean variable `train_on_gpu`. Later, you'll be responsible for making sure that `> * Models, * Model inputs, and * Loss function arguments`

Are moved to GPU, where appropriate.

```

In [16]: """
        DON'T MODIFY ANYTHING IN THIS CELL
        """
        import torch

        # Check for a GPU
        train_on_gpu = torch.cuda.is_available()
        if not train_on_gpu:
            print('No GPU found. Please use a GPU to train your neural network.')
        else:
            print('Training on GPU!')

```

Training on GPU!

---

## 2.5 Discriminator and Generator Losses

Now we need to calculate the losses for both types of adversarial networks.

### 2.5.1 Discriminator Losses

- For the discriminator, the total loss is the sum of the losses for real and fake images,  $d\_loss = d\_real\_loss + d\_fake\_loss$ .
- Remember that we want the discriminator to output 1 for real images and 0 for fake images, so we need to set up the losses to reflect that.

### 2.5.2 Generator Loss

The generator loss will look similar only with flipped labels. The generator's goal is to get the discriminator to *think* its generated images are *real*.

**Exercise: Complete real and fake loss functions** You may choose to use either cross entropy or a least squares error loss to complete the following `real_loss` and `fake_loss` functions.

```
In [38]: def real_loss(D_out):
        '''Calculates how close discriminator outputs are to being real.
           param, D_out: discriminator logits
           return: real loss'''

        batch_size = D_out.size(0)
        # label smoothing
        labels = torch.ones(batch_size) * 0.9
        if train_on_gpu:
            labels = labels.cuda()
        criterion = nn.BCEWithLogitsLoss()
        loss = criterion(D_out.squeeze(), labels)
        return loss

    def fake_loss(D_out):
        '''Calculates how close discriminator outputs are to being fake.
           param, D_out: discriminator logits
           return: fake loss'''
        batch_size = D_out.size(0)
        labels = torch.zeros(batch_size) # fake labels = 0
        if train_on_gpu:
            labels = labels.cuda()
        criterion = nn.BCEWithLogitsLoss()
        # calculate loss
        loss = criterion(D_out.squeeze(), labels)
        return loss
```

## 2.6 Optimizers

**Exercise: Define optimizers for your Discriminator (D) and Generator (G)** Define optimizers for your models with appropriate hyperparameters.

```
In [48]: import torch.optim as optim
        lr = 0.0002
```

```

beta1= 0.5
beta2= 0.99

# Create optimizers for the discriminator D and generator G
d_optimizer = optim.Adam(D.parameters(), lr, [beta1, beta2])
g_optimizer = optim.Adam(G.parameters(), lr, [beta1, beta2])

```

---

## 2.7 Training

Training will involve alternating between training the discriminator and the generator. You'll use your functions `real_loss` and `fake_loss` to help you calculate the discriminator losses.

- You should train the discriminator by alternating on real and fake images
- Then the generator, which tries to trick the discriminator and should have an opposing loss function

**Saving Samples** You've been given some code to print out some loss statistics and save some generated "fake" samples.

**Exercise: Complete the training function** Keep in mind that, if you've moved your models to GPU, you'll also have to move any model inputs to GPU.

```

In [46]: def train(D, G, n_epochs, print_every=50):
    '''Trains adversarial networks for some number of epochs
    param, D: the discriminator network
    param, G: the generator network
    param, n_epochs: number of epochs to train for
    param, print_every: when to print and record the models' losses
    return: D and G losses'''

    # move models to GPU
    if train_on_gpu:
        D.cuda()
        G.cuda()

    # D.train()
    # keep track of loss and generated, "fake" samples
    samples = []
    losses = []

    # Get some fixed data for sampling. These are images that are held
    # constant throughout training, and allow us to inspect the model's performance
    sample_size=16
    fixed_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
    fixed_z = torch.from_numpy(fixed_z).float()
    # move z to GPU if available

```

```

if train_on_gpu:
    fixed_z = fixed_z.cuda()

# epoch training loop
for epoch in range(n_epochs):

    # batch training loop
    for batch_i, (real_images, _) in enumerate(celeba_train_loader):

        batch_size = real_images.size(0)
        real_images = scale(real_images)

        # =====
        #          YOUR CODE HERE: TRAIN THE NETWORKS
        # =====

        # 1. Train the discriminator on real and fake images

        d_optimizer.zero_grad()

        if train_on_gpu:
            real_images = real_images.cuda()

        d_real = D(real_images)
        d_real_loss = real_loss(d_real)

        # Generate fake images
        z = np.random.uniform(-1, 1, size=(batch_size, z_size))
        z = torch.from_numpy(z).float()

        if train_on_gpu:
            z = z.cuda()

        fake_images = G(z)

        d_fake = D(fake_images)
        d_fake_loss = fake_loss(d_fake)

        # add up loss and perform backprop
        d_loss = d_real_loss + d_fake_loss

        d_loss.backward()
        d_optimizer.step()

        # 2. Train the generator with an adversarial loss
        g_optimizer.zero_grad()

        # Generate fake images

```

```

z = np.random.uniform(-1, 1, size=(batch_size, z_size))
z = torch.from_numpy(z).float()

if train_on_gpu:
    z = z.cuda()

fake_images = G(z)
# D.eval()
# Compute the discriminator losses on fake images
# using flipped labels!
g_fake = D(fake_images)
g_loss = real_loss(g_fake) # use real loss to flip labels
# g_loss = torch.mean(g_fake_out**2)
# D.train()
# perform backprop
g_loss.backward()
g_optimizer.step()

# =====
#                     END OF YOUR CODE
# =====

# Print some loss stats
if batch_i % print_every == 0:
    # append discriminator loss and generator loss
    losses.append((d_loss.item(), g_loss.item()))
    # print discriminator and generator loss
    print('Epoch [{:5d}/{:5d}] | d_loss: {:.4f} | g_loss: {:.4f}'.format(
        epoch+1, n_epochs, d_loss.item(), g_loss.item()))

## AFTER EACH EPOCH##
# this code assumes your generator is named G, feel free to change the name
# generate and save sample, fake images
G.eval() # for generating samples
samples_z = G(fixed_z)
samples.append(samples_z)
G.train() # back to training mode
print('-----Epoch [{:5d}/{:5d}]-----'.format(epoch+1, n_epochs))
# Save training generator samples
with open('train_samples.pkl', 'wb') as f:
    pickle.dump(samples, f)

# finally return losses
return losses

```

Set your number of training epochs and train your GAN!

In [49]: # set number of epochs

```
n_epochs = 5
```

```
"""  
DON'T MODIFY ANYTHING IN THIS CELL  
"""
```

```
# call training function  
losses = train(D, G, n_epochs=n_epochs)
```

```
Epoch [ 1/ 5] | d_loss: 1.3303 | g_loss: 1.8101  
Epoch [ 1/ 5] | d_loss: 0.9189 | g_loss: 2.4704  
Epoch [ 1/ 5] | d_loss: 0.9958 | g_loss: 2.9792  
Epoch [ 1/ 5] | d_loss: 0.9461 | g_loss: 2.6381  
Epoch [ 1/ 5] | d_loss: 1.0764 | g_loss: 2.0608  
Epoch [ 1/ 5] | d_loss: 1.0608 | g_loss: 1.3316  
Epoch [ 1/ 5] | d_loss: 0.9782 | g_loss: 1.9480  
Epoch [ 1/ 5] | d_loss: 1.1256 | g_loss: 2.1165  
Epoch [ 1/ 5] | d_loss: 1.1463 | g_loss: 0.9794  
Epoch [ 1/ 5] | d_loss: 0.9937 | g_loss: 2.7013  
Epoch [ 1/ 5] | d_loss: 1.1576 | g_loss: 1.8628  
Epoch [ 1/ 5] | d_loss: 0.9833 | g_loss: 1.4450  
Epoch [ 1/ 5] | d_loss: 1.1488 | g_loss: 1.0338  
Epoch [ 1/ 5] | d_loss: 1.0447 | g_loss: 2.1604  
Epoch [ 1/ 5] | d_loss: 1.2242 | g_loss: 1.6244  
Epoch [ 1/ 5] | d_loss: 1.1264 | g_loss: 1.6144  
Epoch [ 1/ 5] | d_loss: 1.0533 | g_loss: 1.9974  
Epoch [ 1/ 5] | d_loss: 1.1206 | g_loss: 1.4049  
Epoch [ 1/ 5] | d_loss: 1.2711 | g_loss: 1.6288  
Epoch [ 1/ 5] | d_loss: 1.0261 | g_loss: 1.5810  
Epoch [ 1/ 5] | d_loss: 1.2386 | g_loss: 1.3667  
Epoch [ 1/ 5] | d_loss: 1.4252 | g_loss: 2.0231  
Epoch [ 1/ 5] | d_loss: 1.3349 | g_loss: 1.4028  
Epoch [ 1/ 5] | d_loss: 1.1925 | g_loss: 1.4533  
Epoch [ 1/ 5] | d_loss: 1.0897 | g_loss: 1.7333  
Epoch [ 1/ 5] | d_loss: 1.0480 | g_loss: 1.5573  
Epoch [ 1/ 5] | d_loss: 1.1793 | g_loss: 1.8161  
Epoch [ 1/ 5] | d_loss: 1.0713 | g_loss: 1.2107  
Epoch [ 1/ 5] | d_loss: 1.1129 | g_loss: 1.9358  
Epoch [ 1/ 5] | d_loss: 1.0904 | g_loss: 1.1871  
Epoch [ 1/ 5] | d_loss: 1.1940 | g_loss: 1.4800  
Epoch [ 1/ 5] | d_loss: 1.2391 | g_loss: 1.3657  
Epoch [ 1/ 5] | d_loss: 1.1433 | g_loss: 1.5695  
Epoch [ 1/ 5] | d_loss: 1.1558 | g_loss: 1.3800  
Epoch [ 1/ 5] | d_loss: 1.3519 | g_loss: 1.4693  
Epoch [ 1/ 5] | d_loss: 1.0961 | g_loss: 0.8930  
Epoch [ 1/ 5] | d_loss: 1.1069 | g_loss: 1.0004  
Epoch [ 1/ 5] | d_loss: 1.0716 | g_loss: 0.9899  
Epoch [ 1/ 5] | d_loss: 1.1847 | g_loss: 0.9073
```

Epoch [	1/	5]	d_loss: 1.2216	g_loss: 0.9398
Epoch [	1/	5]	d_loss: 1.1167	g_loss: 1.1747
Epoch [	1/	5]	d_loss: 1.4117	g_loss: 0.9760
Epoch [	1/	5]	d_loss: 1.3926	g_loss: 1.3685
Epoch [	1/	5]	d_loss: 1.3819	g_loss: 0.8303
Epoch [	1/	5]	d_loss: 1.2440	g_loss: 1.0139
Epoch [	1/	5]	d_loss: 1.2711	g_loss: 1.3819
Epoch [	1/	5]	d_loss: 1.3498	g_loss: 1.3272
Epoch [	1/	5]	d_loss: 1.0687	g_loss: 1.8989
Epoch [	1/	5]	d_loss: 1.2891	g_loss: 1.5795
Epoch [	1/	5]	d_loss: 1.2074	g_loss: 1.7629
Epoch [	1/	5]	d_loss: 1.1507	g_loss: 0.9320
Epoch [	1/	5]	d_loss: 1.1847	g_loss: 1.2096
Epoch [	1/	5]	d_loss: 1.0529	g_loss: 1.5578
Epoch [	1/	5]	d_loss: 1.1819	g_loss: 1.3015
Epoch [	1/	5]	d_loss: 1.1715	g_loss: 1.4931
-----Epoch [ 1/ 5]-----				
Epoch [	2/	5]	d_loss: 1.1627	g_loss: 1.6272
Epoch [	2/	5]	d_loss: 1.1574	g_loss: 1.1744
Epoch [	2/	5]	d_loss: 0.9652	g_loss: 1.9074
Epoch [	2/	5]	d_loss: 1.1130	g_loss: 0.9996
Epoch [	2/	5]	d_loss: 1.2565	g_loss: 2.2107
Epoch [	2/	5]	d_loss: 1.0739	g_loss: 1.2958
Epoch [	2/	5]	d_loss: 1.1988	g_loss: 1.3033
Epoch [	2/	5]	d_loss: 1.1181	g_loss: 1.1162
Epoch [	2/	5]	d_loss: 1.0106	g_loss: 1.2677
Epoch [	2/	5]	d_loss: 1.1012	g_loss: 1.4630
Epoch [	2/	5]	d_loss: 1.3881	g_loss: 1.1198
Epoch [	2/	5]	d_loss: 1.3192	g_loss: 1.0885
Epoch [	2/	5]	d_loss: 1.0824	g_loss: 1.4317
Epoch [	2/	5]	d_loss: 1.0844	g_loss: 1.1226
Epoch [	2/	5]	d_loss: 1.0748	g_loss: 1.1634
Epoch [	2/	5]	d_loss: 1.0967	g_loss: 1.6271
Epoch [	2/	5]	d_loss: 1.0849	g_loss: 1.6522
Epoch [	2/	5]	d_loss: 1.0246	g_loss: 1.7299
Epoch [	2/	5]	d_loss: 1.0046	g_loss: 1.4066
Epoch [	2/	5]	d_loss: 1.0270	g_loss: 1.1948
Epoch [	2/	5]	d_loss: 1.0424	g_loss: 1.7010
Epoch [	2/	5]	d_loss: 1.0945	g_loss: 1.0000
Epoch [	2/	5]	d_loss: 1.1017	g_loss: 1.6701
Epoch [	2/	5]	d_loss: 1.1835	g_loss: 1.5545
Epoch [	2/	5]	d_loss: 1.0190	g_loss: 1.7677
Epoch [	2/	5]	d_loss: 1.1803	g_loss: 1.1579
Epoch [	2/	5]	d_loss: 1.1818	g_loss: 1.1371
Epoch [	2/	5]	d_loss: 1.1402	g_loss: 1.4203
Epoch [	2/	5]	d_loss: 1.3288	g_loss: 0.7296
Epoch [	2/	5]	d_loss: 1.1618	g_loss: 1.0568
Epoch [	2/	5]	d_loss: 1.1307	g_loss: 0.9331



Epoch [	2/	5]	d_loss: 1.2627	g_loss: 1.4149
Epoch [	2/	5]	d_loss: 1.0511	g_loss: 1.0326
Epoch [	2/	5]	d_loss: 1.0870	g_loss: 0.7921
Epoch [	2/	5]	d_loss: 1.1752	g_loss: 1.0170
Epoch [	2/	5]	d_loss: 0.9872	g_loss: 2.2428
Epoch [	2/	5]	d_loss: 0.9704	g_loss: 1.5596
Epoch [	2/	5]	d_loss: 1.1806	g_loss: 1.2286
Epoch [	2/	5]	d_loss: 1.1301	g_loss: 1.1052
Epoch [	2/	5]	d_loss: 1.2500	g_loss: 0.9430
Epoch [	2/	5]	d_loss: 1.1189	g_loss: 1.3304
Epoch [	2/	5]	d_loss: 1.0920	g_loss: 1.4507
Epoch [	2/	5]	d_loss: 1.2245	g_loss: 1.3676
Epoch [	2/	5]	d_loss: 1.2325	g_loss: 1.1162
Epoch [	2/	5]	d_loss: 1.0051	g_loss: 1.3649
Epoch [	2/	5]	d_loss: 1.0410	g_loss: 1.1180
Epoch [	2/	5]	d_loss: 1.0804	g_loss: 1.0517
Epoch [	2/	5]	d_loss: 1.1395	g_loss: 0.8095
Epoch [	2/	5]	d_loss: 1.0239	g_loss: 1.5515
Epoch [	2/	5]	d_loss: 1.0218	g_loss: 1.3618
Epoch [	2/	5]	d_loss: 1.0925	g_loss: 1.3726
Epoch [	2/	5]	d_loss: 1.0264	g_loss: 1.5255
Epoch [	2/	5]	d_loss: 1.1533	g_loss: 1.3677
Epoch [	2/	5]	d_loss: 1.1058	g_loss: 1.1827
Epoch [	2/	5]	d_loss: 1.0797	g_loss: 1.5106
-----Epoch [ 2/ 5]-----				
Epoch [	3/	5]	d_loss: 1.1000	g_loss: 1.3319
Epoch [	3/	5]	d_loss: 1.0873	g_loss: 1.1500
Epoch [	3/	5]	d_loss: 1.1452	g_loss: 1.7198
Epoch [	3/	5]	d_loss: 1.0845	g_loss: 1.0732
Epoch [	3/	5]	d_loss: 1.0860	g_loss: 1.0945
Epoch [	3/	5]	d_loss: 1.0938	g_loss: 1.6446
Epoch [	3/	5]	d_loss: 1.2805	g_loss: 1.0559
Epoch [	3/	5]	d_loss: 1.1542	g_loss: 2.0289
Epoch [	3/	5]	d_loss: 1.1909	g_loss: 0.9805
Epoch [	3/	5]	d_loss: 1.1098	g_loss: 0.9959
Epoch [	3/	5]	d_loss: 1.0269	g_loss: 0.9256
Epoch [	3/	5]	d_loss: 1.0814	g_loss: 1.2544
Epoch [	3/	5]	d_loss: 0.9906	g_loss: 1.0117
Epoch [	3/	5]	d_loss: 1.0729	g_loss: 1.8364
Epoch [	3/	5]	d_loss: 1.2751	g_loss: 1.2982
Epoch [	3/	5]	d_loss: 0.9977	g_loss: 1.3714
Epoch [	3/	5]	d_loss: 1.0661	g_loss: 1.5213
Epoch [	3/	5]	d_loss: 1.0772	g_loss: 1.0108
Epoch [	3/	5]	d_loss: 1.0085	g_loss: 1.5079
Epoch [	3/	5]	d_loss: 1.1231	g_loss: 1.4835
Epoch [	3/	5]	d_loss: 1.4694	g_loss: 1.7580
Epoch [	3/	5]	d_loss: 1.1405	g_loss: 1.0256
Epoch [	3/	5]	d_loss: 1.2566	g_loss: 1.6881

Epoch [	3/	5]	d_loss: 1.1109	g_loss: 1.4400
Epoch [	3/	5]	d_loss: 1.0676	g_loss: 1.0189
Epoch [	3/	5]	d_loss: 0.9563	g_loss: 1.7508
Epoch [	3/	5]	d_loss: 1.3666	g_loss: 1.5719
Epoch [	3/	5]	d_loss: 1.1240	g_loss: 0.9188
Epoch [	3/	5]	d_loss: 1.2709	g_loss: 2.5072
Epoch [	3/	5]	d_loss: 1.1778	g_loss: 1.3333
Epoch [	3/	5]	d_loss: 0.9441	g_loss: 1.3916
Epoch [	3/	5]	d_loss: 1.1445	g_loss: 1.2752
Epoch [	3/	5]	d_loss: 1.4651	g_loss: 1.0675
Epoch [	3/	5]	d_loss: 1.2772	g_loss: 0.9323
Epoch [	3/	5]	d_loss: 1.1410	g_loss: 1.1289
Epoch [	3/	5]	d_loss: 1.1266	g_loss: 1.1640
Epoch [	3/	5]	d_loss: 0.9976	g_loss: 0.8858
Epoch [	3/	5]	d_loss: 1.0037	g_loss: 2.0342
Epoch [	3/	5]	d_loss: 0.8988	g_loss: 2.1061
Epoch [	3/	5]	d_loss: 1.5124	g_loss: 1.7126
Epoch [	3/	5]	d_loss: 1.0760	g_loss: 1.4591
Epoch [	3/	5]	d_loss: 1.1751	g_loss: 1.0302
Epoch [	3/	5]	d_loss: 1.1298	g_loss: 0.9334
Epoch [	3/	5]	d_loss: 1.0711	g_loss: 1.5438
Epoch [	3/	5]	d_loss: 0.9823	g_loss: 1.4549
Epoch [	3/	5]	d_loss: 1.2090	g_loss: 1.4320
Epoch [	3/	5]	d_loss: 1.0282	g_loss: 1.3394
Epoch [	3/	5]	d_loss: 1.2050	g_loss: 1.4569
Epoch [	3/	5]	d_loss: 1.2131	g_loss: 2.0083
Epoch [	3/	5]	d_loss: 0.9533	g_loss: 1.1056
Epoch [	3/	5]	d_loss: 1.0064	g_loss: 1.4138
Epoch [	3/	5]	d_loss: 1.0988	g_loss: 1.5264
Epoch [	3/	5]	d_loss: 1.0082	g_loss: 1.5590
Epoch [	3/	5]	d_loss: 1.1167	g_loss: 0.9503
Epoch [	3/	5]	d_loss: 1.3031	g_loss: 0.8754
-----Epoch [ 3/ 5]-----				
Epoch [	4/	5]	d_loss: 1.1631	g_loss: 1.0342
Epoch [	4/	5]	d_loss: 1.1896	g_loss: 1.5345
Epoch [	4/	5]	d_loss: 1.1322	g_loss: 1.1642
Epoch [	4/	5]	d_loss: 1.1228	g_loss: 1.2765
Epoch [	4/	5]	d_loss: 1.0777	g_loss: 1.0005
Epoch [	4/	5]	d_loss: 1.0334	g_loss: 1.3469
Epoch [	4/	5]	d_loss: 0.9633	g_loss: 1.3444
Epoch [	4/	5]	d_loss: 1.3342	g_loss: 3.0526
Epoch [	4/	5]	d_loss: 1.1330	g_loss: 1.2455
Epoch [	4/	5]	d_loss: 1.0691	g_loss: 1.4331
Epoch [	4/	5]	d_loss: 1.1526	g_loss: 1.0826
Epoch [	4/	5]	d_loss: 1.1560	g_loss: 1.2532
Epoch [	4/	5]	d_loss: 1.0685	g_loss: 1.2925
Epoch [	4/	5]	d_loss: 1.3204	g_loss: 0.9528
Epoch [	4/	5]	d_loss: 1.1147	g_loss: 1.3582

Epoch [	4/	5]	d_loss: 1.1402	g_loss: 1.3257
Epoch [	4/	5]	d_loss: 1.2576	g_loss: 1.1800
Epoch [	4/	5]	d_loss: 1.2081	g_loss: 1.1648
Epoch [	4/	5]	d_loss: 1.1446	g_loss: 1.1924
Epoch [	4/	5]	d_loss: 1.0679	g_loss: 1.1193
Epoch [	4/	5]	d_loss: 1.6515	g_loss: 0.6358
Epoch [	4/	5]	d_loss: 1.2448	g_loss: 1.4196
Epoch [	4/	5]	d_loss: 1.0860	g_loss: 1.1965
Epoch [	4/	5]	d_loss: 1.1412	g_loss: 1.4594
Epoch [	4/	5]	d_loss: 1.1428	g_loss: 1.8324
Epoch [	4/	5]	d_loss: 1.5442	g_loss: 1.1552
Epoch [	4/	5]	d_loss: 1.0780	g_loss: 1.3112
Epoch [	4/	5]	d_loss: 1.0376	g_loss: 1.4759
Epoch [	4/	5]	d_loss: 1.1290	g_loss: 1.2284
Epoch [	4/	5]	d_loss: 0.9387	g_loss: 1.2377
Epoch [	4/	5]	d_loss: 1.1796	g_loss: 1.0022
Epoch [	4/	5]	d_loss: 1.0793	g_loss: 1.1137
Epoch [	4/	5]	d_loss: 1.1039	g_loss: 0.9580
Epoch [	4/	5]	d_loss: 1.1311	g_loss: 1.2953
Epoch [	4/	5]	d_loss: 1.1277	g_loss: 1.2415
Epoch [	4/	5]	d_loss: 1.1065	g_loss: 1.3449
Epoch [	4/	5]	d_loss: 1.0006	g_loss: 1.3867
Epoch [	4/	5]	d_loss: 1.1558	g_loss: 1.2903
Epoch [	4/	5]	d_loss: 1.2643	g_loss: 1.0820
Epoch [	4/	5]	d_loss: 1.3165	g_loss: 1.1980
Epoch [	4/	5]	d_loss: 1.0459	g_loss: 0.9204
Epoch [	4/	5]	d_loss: 1.0697	g_loss: 0.9093
Epoch [	4/	5]	d_loss: 1.0334	g_loss: 1.0981
Epoch [	4/	5]	d_loss: 1.0242	g_loss: 1.3021
Epoch [	4/	5]	d_loss: 1.0894	g_loss: 1.8243
Epoch [	4/	5]	d_loss: 1.0142	g_loss: 1.2327
Epoch [	4/	5]	d_loss: 0.9746	g_loss: 1.9779
Epoch [	4/	5]	d_loss: 1.0978	g_loss: 1.3948
Epoch [	4/	5]	d_loss: 1.1170	g_loss: 1.5695
Epoch [	4/	5]	d_loss: 1.0449	g_loss: 1.4499
Epoch [	4/	5]	d_loss: 1.0132	g_loss: 1.6129
Epoch [	4/	5]	d_loss: 1.0961	g_loss: 1.3747
Epoch [	4/	5]	d_loss: 1.1106	g_loss: 1.2598
Epoch [	4/	5]	d_loss: 1.0414	g_loss: 1.1297
Epoch [	4/	5]	d_loss: 1.0223	g_loss: 0.9315
-----Epoch [ 4/ 5]-----				
Epoch [	5/	5]	d_loss: 1.5800	g_loss: 1.0662
Epoch [	5/	5]	d_loss: 1.0986	g_loss: 1.5570
Epoch [	5/	5]	d_loss: 1.3748	g_loss: 1.4584
Epoch [	5/	5]	d_loss: 1.0725	g_loss: 1.3480
Epoch [	5/	5]	d_loss: 1.3001	g_loss: 1.0941
Epoch [	5/	5]	d_loss: 1.6287	g_loss: 0.9990
Epoch [	5/	5]	d_loss: 1.0766	g_loss: 1.0211

Epoch [	5/	5]	d_loss: 1.0724	g_loss: 1.4202
Epoch [	5/	5]	d_loss: 1.2767	g_loss: 1.2965
Epoch [	5/	5]	d_loss: 1.0913	g_loss: 1.7250
Epoch [	5/	5]	d_loss: 1.3201	g_loss: 1.3872
Epoch [	5/	5]	d_loss: 1.1233	g_loss: 1.1420
Epoch [	5/	5]	d_loss: 1.0417	g_loss: 1.4755
Epoch [	5/	5]	d_loss: 1.1008	g_loss: 1.3268
Epoch [	5/	5]	d_loss: 1.0311	g_loss: 1.7044
Epoch [	5/	5]	d_loss: 1.1262	g_loss: 0.9980
Epoch [	5/	5]	d_loss: 1.0637	g_loss: 1.2031
Epoch [	5/	5]	d_loss: 1.2018	g_loss: 1.2482
Epoch [	5/	5]	d_loss: 1.0289	g_loss: 1.4648
Epoch [	5/	5]	d_loss: 1.4631	g_loss: 1.8113
Epoch [	5/	5]	d_loss: 1.2087	g_loss: 1.2641
Epoch [	5/	5]	d_loss: 1.0632	g_loss: 1.2208
Epoch [	5/	5]	d_loss: 1.0690	g_loss: 1.1960
Epoch [	5/	5]	d_loss: 1.2447	g_loss: 1.6455
Epoch [	5/	5]	d_loss: 1.0945	g_loss: 0.9510
Epoch [	5/	5]	d_loss: 1.0486	g_loss: 1.2811
Epoch [	5/	5]	d_loss: 1.1892	g_loss: 1.4895
Epoch [	5/	5]	d_loss: 1.1955	g_loss: 1.6984
Epoch [	5/	5]	d_loss: 1.0672	g_loss: 1.5833
Epoch [	5/	5]	d_loss: 1.0685	g_loss: 1.6787
Epoch [	5/	5]	d_loss: 1.0731	g_loss: 1.4402
Epoch [	5/	5]	d_loss: 1.1193	g_loss: 1.3584
Epoch [	5/	5]	d_loss: 1.1736	g_loss: 0.8729
Epoch [	5/	5]	d_loss: 1.1918	g_loss: 1.8786
Epoch [	5/	5]	d_loss: 1.0325	g_loss: 1.3899
Epoch [	5/	5]	d_loss: 1.1024	g_loss: 0.9886
Epoch [	5/	5]	d_loss: 0.9862	g_loss: 1.5016
Epoch [	5/	5]	d_loss: 1.2605	g_loss: 1.6372
Epoch [	5/	5]	d_loss: 1.0174	g_loss: 2.9381
Epoch [	5/	5]	d_loss: 1.1860	g_loss: 1.4207
Epoch [	5/	5]	d_loss: 1.2058	g_loss: 1.0968
Epoch [	5/	5]	d_loss: 0.9666	g_loss: 1.4738
Epoch [	5/	5]	d_loss: 1.0771	g_loss: 1.6436
Epoch [	5/	5]	d_loss: 1.0786	g_loss: 1.5552
Epoch [	5/	5]	d_loss: 0.8524	g_loss: 1.4676
Epoch [	5/	5]	d_loss: 1.0555	g_loss: 1.4262
Epoch [	5/	5]	d_loss: 0.9890	g_loss: 1.1580
Epoch [	5/	5]	d_loss: 1.0852	g_loss: 1.1590
Epoch [	5/	5]	d_loss: 1.0275	g_loss: 1.5473
Epoch [	5/	5]	d_loss: 1.0663	g_loss: 1.3242
Epoch [	5/	5]	d_loss: 0.9070	g_loss: 1.4726
Epoch [	5/	5]	d_loss: 1.3565	g_loss: 0.8893
Epoch [	5/	5]	d_loss: 1.1367	g_loss: 1.4613
Epoch [	5/	5]	d_loss: 1.2253	g_loss: 2.0900
Epoch [	5/	5]	d_loss: 1.0446	g_loss: 0.8971


-----Epoch [ 5/ 5]-----

## 2.8 Training loss

Plot the training losses for the generator and discriminator, recorded after each epoch.

```
In [50]: fig, ax = plt.subplots()
         losses = np.array(losses)
         plt.plot(losses.T[0], label='Discriminator', alpha=0.5)
         plt.plot(losses.T[1], label='Generator', alpha=0.5)
         plt.title("Training Losses")
         plt.legend()
```

```
Out[50]: <matplotlib.legend.Legend at 0x7f01f018ad30>
```



## 2.9 Generator samples from training

View samples of images from the generator, and answer a question about the strengths and weaknesses of your trained models.

```
In [51]: # helper function for viewing a list of passed in sample images
         def view_samples(epoch, samples):
             fig, axes = plt.subplots(figsize=(16,4), nrows=2, ncols=8, sharey=True, sharex=True)
             for ax, img in zip(axes.flatten(), samples[epoch]):
```

```

img = img.detach().cpu().numpy()
img = np.transpose(img, (1, 2, 0))
img = ((img + 1)*255 / (2)).astype(np.uint8)
ax.xaxis.set_visible(False)
ax.yaxis.set_visible(False)
im = ax.imshow(img.reshape((32,32,3)))

```

```


In [52]: # Load samples from generator, taken while training
         with open('train_samples.pkl', 'rb') as f:
             samples = pkl.load(f)

```

```

In [53]: _ = view_samples(-1, samples)

```



### 2.9.1 Question: What do you notice about your generated samples and how might you improve this model?

When you answer this question, consider the following factors: \* The dataset is biased; it is made of "celebrity" faces that are mostly white \* Model size; larger models have the opportunity to learn more features in a data feature space \* Optimization strategy; optimizers and number of epochs affect your final result

**Answer:** - I tried various hyper parameters like - lr=0.0005, beta1=0.2 - lr=0.0002, beta1=0.5 The 2nd option worked but gave more g\_loss. So I analyzed some blogs and papers and got the suggestion that I should change batchsize, so I changed it from 16 to 32. Then the generating loss reduced gradually. - As I analyzed the generated samples, I wanted my images to be more clear with discriminator and generator loss as close as possible. Also wanted some diversity in ages and skin color. - I think of making the following improvements: - based on suggestions, provided by mentor [Knowledge center](#). I think of increasing the image size provided by the generator to 128x128, then resizing it to 32x32 and providing it to discriminator. - Also using a different dataset of cartoons or faces of all diversity to train images would help like one by [IBM](#) - Maybe using some other method for smoothing the value for loss calculation.

### 2.9.2 Submitting This Project

When submitting this project, make sure to run all the cells before saving the notebook. Save the notebook file as "dlnd\_face\_generation.ipynb" and save it as a HTML file under "File" -> "Download as". Include the "problem\_unittests.py" files in your submission.